

Computer Architecture 2024 Fall Lab 2

1. Problem Description

In this homework, you are going to extend your Lab 1 to a pipelined CPU. This CPU has 32 registers and 1KB instruction memory. It should take 32-bit binary codes as input and should do the corresponding RISC-V instructions, saving the result of arithmetic operations into the corresponding registers. Besides the instruction specified in lab 1, you have to support `lw`, `sw`, `beq` in this Lab additionally. We will examine the correctness of your implementation by dumping the value of each register and data memory after each cycle.

1.1. Pipeline Registers

To support pipeline execution, the first step is adding pipeline registers to the CPU. Pipeline registers store control signals and data from the last step and isolate each step. Note that you have to use the non-blocking assignment in your pipeline register to get the correct result. And you have to properly initialize reg values in your pipeline registers.

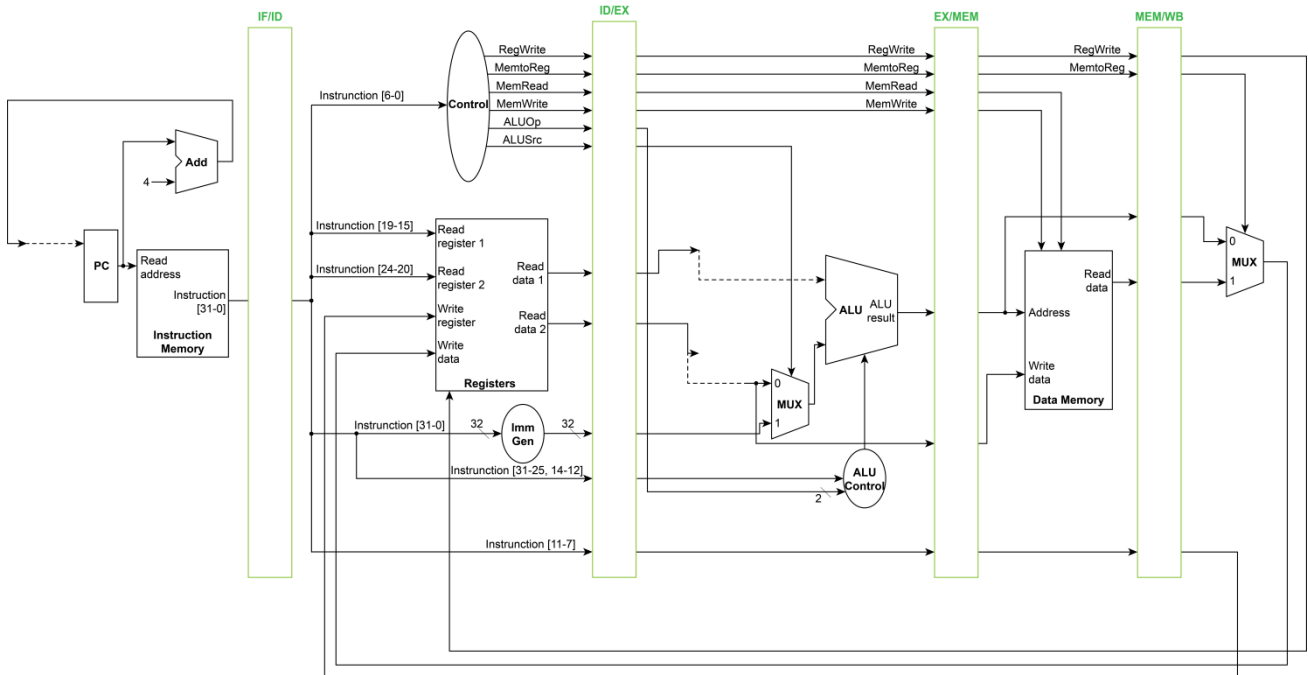


Figure 2 Datapath after adding pipeline registers (colored in green)

1.2. Data Hazard and Forwarding Control

The first issue after adding pipeline execution is to handle data hazards. For example, if we execution “add t3, t1, t2” and “add t4, t1, t3” sequentially, t3 will not be written back until the 5th cycle. But the second instruction is brought into the execution stage at the 4th cycle, causing a different result from single-cycle implementation. To properly handle such data hazards, we can forward the ALU result of the first instruction to the execution stage of the second instruction.

As a result, we shall have a “forwarding unit” exploring whether we should forward data from MEM stage or WB stage to EX stage. If the forwarding unit finds that the data in the later stage will be written back to the register that is taken in the EX stage, the data should be forwarded from the later stage to ensure correct arithmetic operation. To be more specific, Table 1, Listing 1, and Listing 2 are the recommended method to resolve data hazards from the textbook.

Note that the forwarding unit is placed in EX stage, implying that we only forward to EX stage. You don’t have to handle the cases that forwarding to ID stage is necessary to keep the correctness. For example,

```
add x5, x6, x7
beq x5, x4, BRANCH
```

instruction sequences like this will not be involved in our evaluation.

MUX	value	Source	Explanation
ForwardA	00	ID/EX	The first ALU operand comes from the register files
	10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
	01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB	00	ID/EX	The second ALU operand comes from the register files
	10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
	01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

Table 1 Forwarding Control

```

if (EX/MEM.RegWrite
and (EX/MEM.RegisterRd != 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRs1)) ForwardA = 10

if (EX/MEM.RegWrite
and (EX/MEM.RegisterRd != 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRs2)) ForwardB = 10

```

Listing 1 EX hazard

```

if (MEM/WB.RegWrite
and (MEM/WB.RegisterRd != 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd != 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRs1))
and (MEM/WB.RegisterRd == ID/EX.RegisterRs1)) ForwardA = 01

if (MEM/WB.RegWrite
and (MEM/WB.RegisterRd != 0)
and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd != 0)
and (EX/MEM.RegisterRd == ID/EX.RegisterRs2))

```

```
and (MEM/WB.RegisterRd == ID/EX.RegisterRs2)) ForwardB = 01
```

Listing 2 MEM hazards

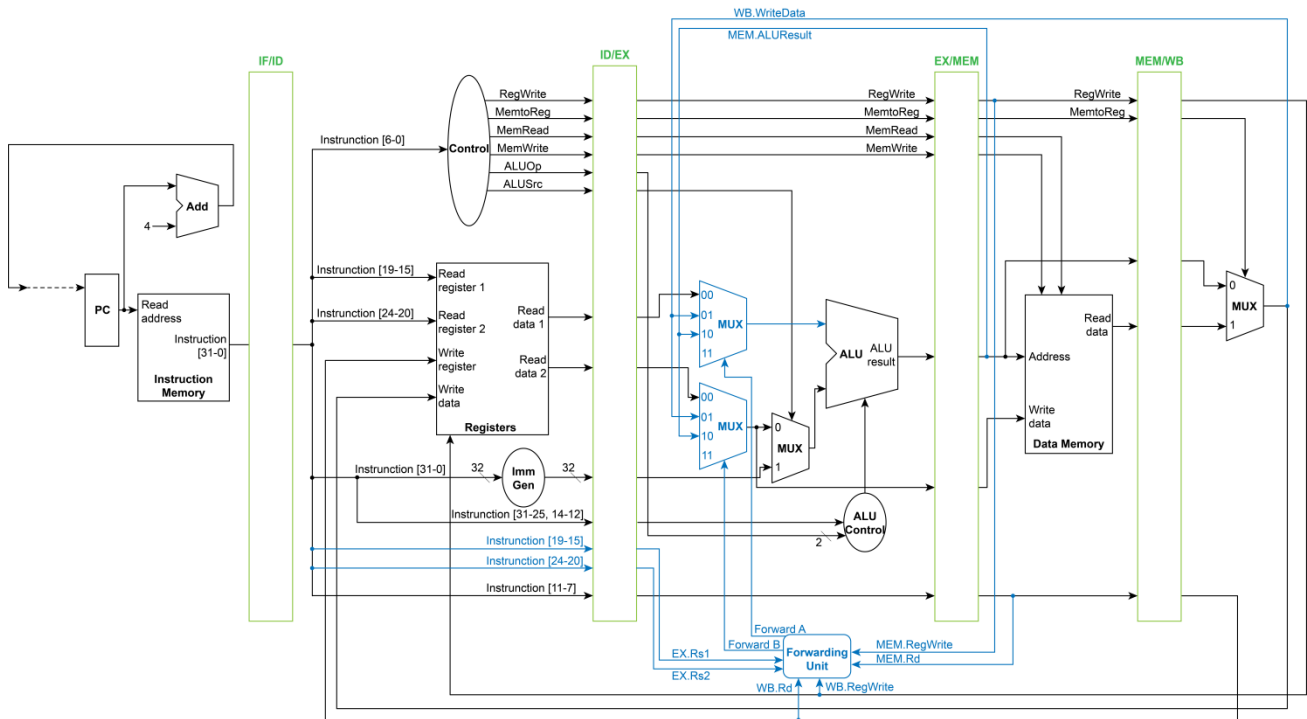


Figure 3 Data path after adding Forwarding Control (colored in blue)

1.3.Hazard Detection, Stall, and Flush

Besides data hazards caused by arithmetic operations, which can be resolved by forwarding, data hazards caused by “load” cannot be resolved simply by forwarding and requires stall. Another major difference between lab 1 and this lab is that we have to support branch instruction, which causes control hazards for the wrong prediction.

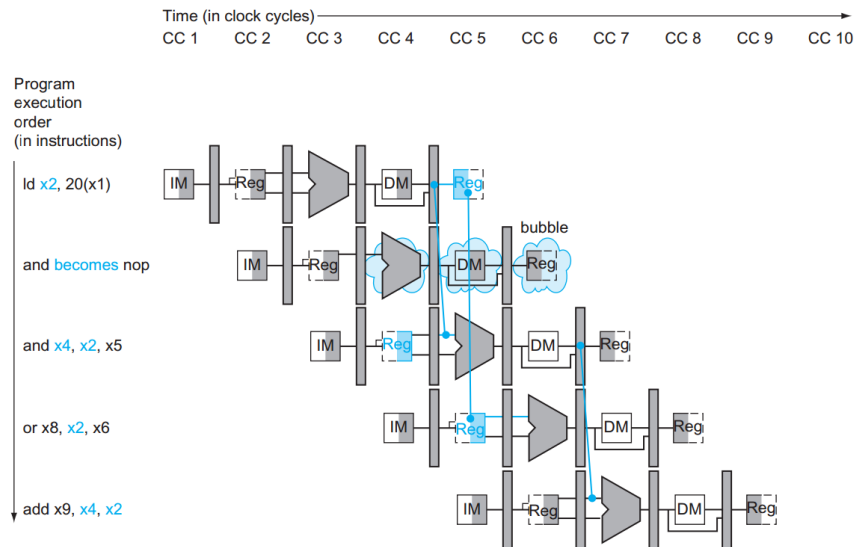


FIGURE 4.57 The way stalls are really inserted into the pipeline. A bubble is inserted beginning in clock cycle 4, by changing the `and` instruction to a `nop`. Note that the `and` instruction is really fetched and decoded in clock cycles 2 and 3, but its EX stage is delayed until clock cycle 5 (versus the unstalled position in clock cycle 4). Likewise, the `or` instruction is fetched in clock cycle 3, but its ID stage is delayed until clock cycle 5 (versus the unstalled clock cycle 4 position). After insertion of the bubble, all the dependencies go forward in time and no further hazards occur.

Figure 4 An example to stall from textbook

As a result, we have to implement a “hazard detection unit” to detect whether to stall the pipeline or to flush when a control hazard happens. The hazard detection unit detects whether the `rd` in EX stage is the same as `rs1` or `rs2` in ID stage. If so, adding a `nop` (no operation) to the pipeline to resolve data hazard.

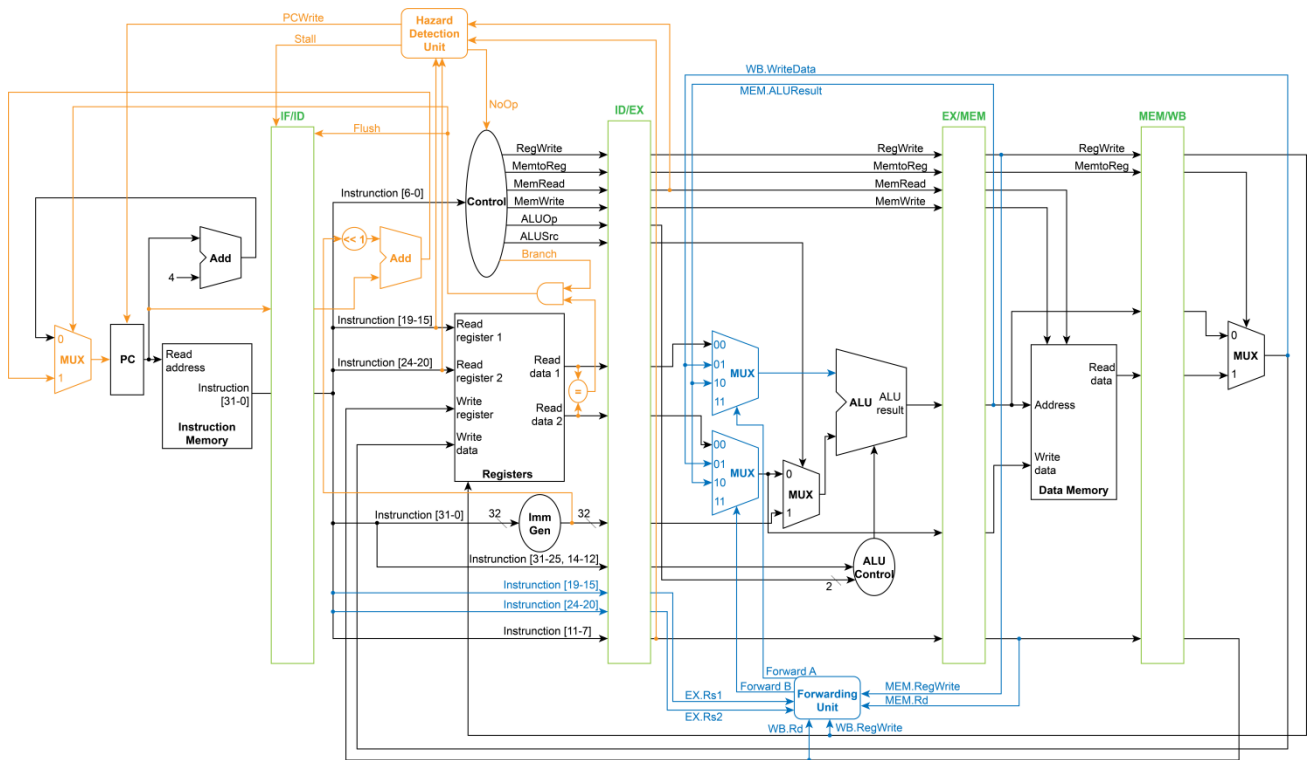


Figure 5 Final Datapath after Adding Hazard Detection Unit (colored in orange)

Note that to mitigate the impact of the branch instruction, we implement the branch decision at ID stage without using ALU, which is a recommended implementation by the textbook. You don't have to handle the forwarding to ID stage. We will avoid the test cases that require forwarding to ID stage (example at the end of 1.3.).

1.4. Instructions

Besides instructions specified in lab 1, you have to support additional 3 instructions, `lw`, `sw`, `beq`. Their machine code is as follows.

funct7	rs2	rs1	funct3	rd	opcode	function
0000000	rs2	rs1	111	rd	0110011	and
0000000	rs2	rs1	100	rd	0110011	xor
0000000	rs2	rs1	001	rd	0110011	sll
0000000	rs2	rs1	000	rd	0110011	add
0100000	rs2	rs1	000	rd	0110011	sub
0000001	rs2	rs1	000	rd	0110011	mul
imm[11:0]		rs1	000	rd	0010011	addi
0100000	imm[4:0]	rs1	101	rd	0010011	srai
imm[11:0]		rs1	010	rd	0000011	lw
imm[11:5]	rs2	rs1	010	imm[4:0]	0100011	sw
imm[12,10:5]	rs2	rs1	000	imm[4:1,11]	1100011	beq

1.5. Input / Output Format

Besides the modules listed above, you are also provided "testbench.v" and "instruction.txt". After you finish your modules and CPU, you should compile all of them including "testbench.v". A recommended compilation command would be

```
$ iverilog *.v -o CPU.out
```

Then by default, your CPU loads "instruction.txt", which should be placed in the same directory as CPU.out, into the instruction memory. This part is written in "testbench.v". You don't have to change it. "instruction.txt" is a plain text file that consists of 32 bits (ASCII 0 or 1) per line, representing one instruction per line. For example, the first 3 lines in "instruction.txt" are

```
0000000_000000_000000_000_01000_0110011 //add $t0,$0,$0
```

```
000000001010_00000_000_01001_0010011 //addi $t1,$0,10  
000000001101_00000_000_01010_0010011 //addi $t2,$0,13
```

Note that underlines and texts after “//” (i.e. comments) are neglected. They are inserted simply for human readability. Therefore, the CPU should take

“00000000000000000000000010000110011” and execute it in the first cycle, then
“00000000101000000000010010010011” in the second cycle, and
“00000000110100000000010100010011” in the third, and so on.

Also, if you include unchanged “testbench.v” into the compilation, the program will generate a plain text file named “output.txt”, which dumps values of all registers and data memory at each cycle after execution. The file is self-explainable.

1.6. Modules You Need to Add or Modify

1.6.1. Control and ALU_Control

Because your CPU has to support load/store instructions in this lab, which are not involved in lab 1, you have to add some additional control signals in the Control and ALU_Control module.

1.6.2. Pipeline Registers

As is introduced in section 1.2., you have to implement 4 pipeline registers to isolate 5 pipeline stages, and passing essential information to the next stage.

1.6.3. Forwarding Unit

As is introduced in section 1.3., you need a forwarding unit and two multiplexers to control forwarding from MEM stage or WB stage.

1.6.4. Hazard Detection Unit

As is introduced in section 1.4., you need a hazard detection unit to handle the necessary stall and nop (no operation).

1.6.5. Others

You can add more modules than listed above if you want. Figure 5 is simply a recommended data path for you to refer to. You are free to change some details as long as your CPU can perform correctly.

1.6.6. CPU

Use structure modeling to connect the input and output of modules following the data path in Figure 5.

1.7. Module Naming Restrictions

To simplify grading, ensure that testbench.v can access your modules without any modifications. That is, name the modules exactly as specified in testbench.v. For example, use names like CPU.Hazard_Detection.Stall_o, CPU.Control.Branch_o, and CPU.ID_FlushIF.

1.8. Reminder

Lab 3 will be strongly related to this homework. Please make sure you can fully understand how to write this homework; otherwise, you may encounter difficulties in your next lab. Plagiarism is strongly prohibited.

2. Report

2.1. Modules Explanation

You should briefly explain how the modules you implement work in the report. You have to explain them in human-readable sentences. Either English or Chinese is welcome, but no Verilog. Explaining Verilog modules in Verilog is nonsense. Simply pasting your codes into the report with no or little explanation will get zero points for the report. You have to write more detail than Section 1.6.

Take “PC.v” as an example, an acceptable report would be:

PC module reads clock signals, reset bit, start bit, and next cycle PC as input, and outputs the PC of the current cycle. This module changes its internal register “pc_o” at the positive edge of the clock signal. When the reset signal is set, PC is reset to 0. And PC will only be updated by next PC when the start bit is on.

And following report will get zero points.

The inputs of PC are clk_i, rst_i, start_i, pc_i, and output pc_o.

It works as follows:

```
always@(posedge clk_i or negedge rst_i) begin
    if(rst_i) begin
        pc_o <= 32'b0;
    end
    else begin
        if(start_i)
            pc_o <= pc_i;
        else
            pc_o <= pc_o;
    end
end
```

2.2. Difficulties Encountered and Solutions in This Lab

Write down the difficulties if any you encountered in doing this lab, and the final solution to them.

3. Environment

Requirement: Docker

After you unzip the supplied file, we will get a directory in the following structure

- Lab2/
 - Lab2_spec.pdf
 - Lab2/
 - ◆ Makefile
 - ◆ code/
 - src/
 - <Implement your CPU here>
 - supplied/
 - <supplied codes, don't modify them>
 - ◆ Docker-compose.yml
 - ◆ Dockerfile
 - ◆ Judge.yaml
 - ◆ Testcases/
 - <sample testcases>
 - ◆ Log/
 - <logs>

You should modify the codes in “code/src” only and do not modify anything in “code/supplied”.

The sample testcases are listed in the directory “testcases”.

After implemented your CPU, use `$make` or `$docker-compose up` to execute your code and get the log and report. Docker-compose will build up the environment for you.

You can find the log, text and waveform, in the directory “log” after each execution.

If you don't have docker in your environment. The alternative is using the following commands

```
$cp testcases/instruction_n.txt instruction.txt
```

```
$iverilog -o cpu code/src/*.v code/supplied/*.v
```

```
$vvp cpu
```

And compare the file “output.txt” with the corresponding output_n.txt

We recommend develop under the provided environment, otherwise, please make sure your code run correctly under **ubuntu 22.04** with **iverilog=11.0-1.1**. **You will get 0 point if your code cannot be executed correctly on our environment.**

4. Submission Rules

Copy all your Verilog codes in “code/src” into a directory named “src”, then put “src” and your report (should be named in format “studentID_lab2_report.pdf”, e.g. “b11902000_lab2_report.pdf”) into a directory named “studentID_lab2”. **Please note that you don’t have to include any other files provided by the TA, such as testbench.v, Instruction_Memory.v, Registers.v, PC.v, instruction.txt, and output.txt, in your submission.** Finally, zip this directory, and upload this zip file onto NTU COOL before **11/26/2024 (Tue.) 23:59**.

In short, we should see a single directory like following structure after we type

```
$ unzip b11902000_lab2.zip
```

In Linux terminal:

- b11902000_lab2/
 - b11902000_lab2/src
 - b11902000_lab2/src/CPU.v
 - b11902000_lab2/src/ALU.v
 - ...
 - b11902000_lab2/b11902000_lab2_report.pdf

Make sure you didn’t includes the following files in your submission: Instruction_Memory.v, PC.v, Registers.v, testbench.v, instruction.txt, output.txt.

5. Evaluation Criteria

5.1. Programming Part (80%)

- We will have a demonstration session on both labs. You have to come up to briefly explain how your program works to get the credits of the programming part.

5.2. Report (20%)

5.3. Other

Minor mistakes (examples below) causing compilation error: -10 pts

- wrong usage of “`include”
- submitting unnecessary files (those provided by TA except CPU.v and testbench.v)
- other mistakes that can be fixed within 5 lines

Major mistakes causing compilation error: 0 pts on programming part

- No show up at demonstration: 0 pts on programming part

Wrong directory format: -10 pts

Wrong I/O paths: -10 pts

Late submission: -10 pts per day

Plagiarism: 0 pts.

email to ntu.eclab.ca.ta@gmail.com if you have any questions.